

Parallelization in loo (current state)

Stan GSOC 2026

What is parallelization?

Serial mirai execution

Currently, parallelization in loo goes through [three paths](#) [1]:

1. Serial execution using `lapply()`
2. Forking through `parallel::mclapply()`
3. Fresh processes through `parallel::parLapply()`

The serial execution path is taken when the `cores` argument is equal to 1, and forking is only available on Linux and MacOS. Since `mirai` [2] neatly subsumes paths 2 and 3, a question we had was what would the performance hit be if we also removed 1—that is, if a user requested serial execution, what do we lose by using `mirai` with a single core.

```
log_ratios <- -loo::example_loglik_matrix()
S <- nrow(log_ratios)
N <- ncol(log_ratios)
tail_len <- loo::n_pareto(r_eff = rep(1, N), S = S)

idx <- seq_len(N)

work <- function(i) {
  loo::do_psis_i(
    log_ratios_i = log_ratios[, i],
    tail_len_i = tail_len[i]
  )$pareto_k
}

mirai::daemons(1)
```

```

results <- bench::mark(
  vapply = vapply(id, work, numeric(1)),
  lapply = lapply(id, work),
  mirai = mirai::mirai_map(
    id,
    work,
    log_ratios = log_ratios,
    tail_len = tail_len
  )[],
  vapply_to_list = vapply(id, work, numeric(1)) |> as.list(),
  purrr_map = purrr::map_dbl(id, work),
  check = same_numeric_values,
  iterations = 200
)

```

Note that `mirai_map()` and `lapply()` results are lists. `same_numeric_values()`'s definition is elided from output but simply checks that results are identical without counting unlisting as part of the benchmark (unlike `vapply_to_list`).

```
results |> summary()
```

```

# A tibble: 5 x 6
  expression      min median `itr/sec` mem_alloc `gc/sec`
  <bch:expr>    <bch:tm> <bch:tm>      <dbl>   <bch:byt>    <dbl>
1 vapply         9.42ms  10.5ms     93.4    3.76MB     15.2
2 lapply         9.45ms   11ms     90.5    3.44MB     15.4
3 mirai        22.43ms  25.9ms     38.2   736.67KB     0.581
4 vapply_to_list  9.59ms  11.5ms     85.2    3.44MB     14.5
5 purrr_map      9.73ms  11.5ms     85.0     4.2MB     14.4

```

```
results |> summary(relative = TRUE)
```

```

# A tibble: 5 x 6
  expression      min median `itr/sec` mem_alloc `gc/sec`
  <bch:expr>    <dbl>  <dbl>      <dbl>    <dbl>    <dbl>
1 vapply         1      1      2.45     5.23     26.2
2 lapply        1.00   1.05     2.37     4.78     26.4
3 mirai        2.38   2.46      1      1      1
4 vapply_to_list 1.02   1.09     2.23     4.78     24.9
5 purrr_map     1.03   1.10     2.23     5.84     24.8

```

NB: the `mem_alloc` column should not be interpreted as total memory usage for `mirai::bench::mark()` records R heap allocations via `utils::Rprofmem()`, which is not able to track daemons' memory usage [2], [3], [4].

`mirai_map()` with a single daemon has a median runtime of ~26ms, roughly 2.5x slower than the fastest in-process approach. Of course, in absolute terms the difference is negligible, adding only about 15ms. The memory usage of all in-process approaches are relatively close to one another. `vapply()`'s small memory overhead may probably be attributed to the extra work it does validating types and lengths and such [4].

In short, it is probably a good idea to directly test the effects of serial `mirai` execution in `loo`, but it is likely that we will find that we should keep the unique single core execution path.

Where is parallelization currently in use?

- `function()`
 - short description
- `function()`
 - short description

Where could we benefit from introducing parallelization?

To investigate this question we use two approaches:

- 1) Inspecting the code base and figuring out at which points parallelization might be helpful (“analytic search”)
- 2) Running performance tests on the `loo` pipeline, investigating where the bottlenecks are, and finding out whether we can improve things with parallelization.

There are some functions in `loo` which could benefit from parallelization, potentially to great effect, as these functions appear to do “embarrassingly parallel” work—meaning that they do independent tasks which seem amenable to speedups by executing those tasks concurrently. These as prime targets for parallelization and somewhat easily identified, we do so by inspection, as covered in the next section.

Analytic search

For each function/place in code where you think parallelization might help, do:

- * Introduction
- * short description of function
- * why do you think that parallelization can improve things?
- * Experiment:
- * implement parallelization
- * record time (define which time you are using) for parallel and unparallel version
- * report results (how much speed up we can get?)
- * Short conclusion

We can see in `elpd_loo_approximation()` that we don't push `cores` through to `loo.function()` so right now, this function only parallelizes if users set the `mc.cores` options as `loo.function()` will pick this up when its called without specifying `cores`. This is a super simple fix and could speed up some code, depending on how users tend to set their `cores`.

Later on in the same function, we can see that parallelization had been thought of but not implemented yet for the delta method approximations in that function (`waic_grad`, `waic_grad_marginal`, `waic_hess`). Each of these loop row-wise over data and shouldn't be too difficult to parallelize. More careful thought must go into determining if parallelization is even worth it, and more careful scrutiny of the callsites of `elpd_loo_approximation()` to ensure it doesn't accidentally induce nested parallelization (see Section)—though at first blush this seems unlikely.

In `waic.function()` `loo` simply doesn't parallelize the work, relying only on a base `lapply()`. This seems like a pretty easy fix, and there doesn't seem to be any reason to not use parallelization. Again, should be careful and benchmark to ensure we aren't accidentally eating performance or RAM, but this could be a large, free win.

Lastly, in `ap_psis.array()`, which mostly just converts array arguments to matrices before passing the buck to the matrix method, we can note that `cores` gets hardcoded to 1 when calling `ap_psis.matrix()`. There doesn't seem to be any real reason to do this, eating the input `cores` like in the first example; unlike there though, we pass `cores = 1`. This is not a huge issue because users probably don't use this function, since most times they would be working directly with matrixes. Further evidencing this hypothesis is [this line](#) which should error out since it references a nonexistent `r_eff`. That obviously needs to be fixed, and this function tested.

Performance testing

- short description of approach for running performance test
- run performance test
- report results
- identify bottlenecks
- provide first ideas whether any of the identified bottlenecks can benefit from parallelization

Vectorization

When browsing the codebase to find places where parallelization could be helpful (see Section), we found two potential speedups—in these places, we may be able to benefit from *vectorizing* functions instead of parallelizing a loop. Vectorization is advantageous in these spots as we would be rewriting R loops as optimized matrix operations [5, Secs. 24.5–24.7].

Matrix algebra is a general example of vectorisation. There loops are executed by highly tuned external libraries like BLAS. If you can figure out a way to use matrix algebra to solve your problem, you’ll often get a very fast solution.

Of course, for these potential improvements, we should first write tests for correctness, then write benchmarks in the PR.

The more important of the two is in `pseudobma_weights()`, when performing a Bayesian bootstrap. This function is the backend for `loo_model_weights()`, and is used to average models using psuedo-Bayesian model averaging, optionally using a Bayesian bootstrap (which is referred to as psudo-BMA+). This pseudo-BMA+ path is the one we care about as it has the elidable loop. We currently loop over Bayesian bootstrap iterations where it seems quite plausible to swap to matrix operations. This could potentially be a large speedup as the loop is large, being of length `BB_n`, which defaults to 1,000 replications. There don’t appear to be any reallocations in the loop, so the speedup will probably be dominated by the efficiency of matrix ops [5, Secs. 5.3.1, 24.6]. This is tracked in #XXX.

Similarly, though less interesting, is a gradient calculation in `stacking_weights()`, just above the previous function. `stacking_weights()` is similarly used for combining models using stacking, i.e., fitting a linear model to weigh the contending models. The gradient function is used for solving for those weights. However, the speedup here is less interesting because we are looping over `K`, the number of models being compared, which is usually small. However, we may as well take the (potential) free performance. This is tracked in #XXX.

These two loops seem like low hanging fruit and well worth the time to implement.

Nested parallelization

When updating `loo` we wish to avoid “nested parallelization”, i.e. attempting to compute something in parallel, when already in a parallel process. Take this nonfunctioning `mirai` code as an example:

```
# parallel over xs
mirai_map(
  xs,
  # parallel over x
```

```
function(x) mirai_map(x, mean)[]
)
```

We first are doing something parallelly over values of `xs`, but then we try to work in parallel again for each value from `xs`! The main problem here is oversubscription: we spawn `n` outer workers, let say, then for each we spawn another `m` inner workers. In the package context though, this implementation would be far more complex, and would probably be hidden completely from users. This becomes a problem when a user sets `n`, the outer parallel workers, to `# cores - 1`¹, as the actual number of parallel tasks could be as large as $n * m$. Oversubscription is bad because, in this context, we are mostly running CPU heavy tasks where each process takes time and needs its own core. Oversubscription is promising more parallelization than what the hardware can provide—note that parallelization usually has some fixed startup costs as well (see Section), so we would probably lose performance. In `loo` specifically, we must be careful when parallelizing functions to ensure we do not accidentally have nested parallelization.

Where do we currently find nested parallelization in `loo`?

- short description of function
- do you see any challenges with the identified functions?
- should nested parallelisation run parallel or sequential?
 - provide experiments to stress your statement

Potential problems with packages that make use of `loo` such that nested parallelization can occur?

However, we know that other packages also call `loo` functions, which is another potential vector for nested parallelization. To this end, it seems like `brms` is the only (Stan) package where this could even be vaguely possible, but doing so seems to necessitate wilful parallelization not provided natively by `brms`. This merits further investigation, but it seems like we can (softly) table concerns regarding external nested parallelization for now.

Make a list of packages of which you have thought of: * package name * Any issues that you see? * (make a short note even if you don't see issues)

¹This is a common heuristic for number of threads/daemons/workers/etc. to use because it leaves 1 core free to do normal computer things, lest the system become unresponsive. Note this doesn't account for memory usage at all—if you have 8n memory but send out n tasks which each use “2” memory, you will probably run out of memory (see Section).

Appendix

```
sessioninfo::session_info(  
  pkgs = c("loo", "mirai", "bench", "purrr"),  
  info = c("platform", "packages"),  
  dependencies = FALSE  
)
```

```
Warning in system2("quarto", "-V", stdout = TRUE, env = paste0("TMPDIR=", :  
running command '"quarto"  
TMPDIR=C:/Users/visru/AppData/Local/Temp/RtmpyImYGr/file2c143505670b -V' had  
status 1
```

- Session info -----

setting	value
version	R version 4.6.0 (2026-04-24 ucrt)
os	Windows 11 x64 (build 26200)
system	x86_64, mingw32
ui	RTerm
language	(EN)
collate	English_United States.utf8
ctype	English_United States.utf8
tz	America/Los_Angeles
date	2026-06-01
pandoc	3.8.3 @ c:\\Program Files\\Positron\\resources\\app\\quarto\\bin\\tools/ (via rmar
quarto	NA @ C:\\Users\\visru\\AppData\\Local\\Programs\\Quarto\\bin\\quarto.exe

- Packages -----

package	*	version	date (UTC)	lib	source
bench		1.1.4	2025-01-16	[1]	RSPM (R 4.6.0)
loo		2.9.0	2025-12-23	[1]	RSPM (R 4.6.0)
mirai		2.7.0	2026-05-08	[1]	RSPM (R 4.6.0)
purrr		1.2.2	2026-04-10	[1]	RSPM (R 4.6.0)

[1] C:/Users/visru/AppData/Local/R/win-library/4.6

[2] C:/Program Files/R/R-4.6.0/library

Memory

Memory usage and loo

As briefly discussed in [Nested Parallelization](#), parallelization may lead to running out of memory (OOM). Copying the construction from that note: if you have, say, 8GB memory but spin up 16 tasks which each use 1GB, you will run out of memory. The math here is quite clear, though it is worst case and is complicated by the profile of the program—if that 1GB is peak usage and it normally doesn’t go so high, and the peaks don’t overlap, then it is possible for you to squeeze past without 16 gigs free. Note, however, that this 1GB usage per parallel worker is assuming that there is no memory that can be shared. If objects are reused across workers, we could use something like `mori` [6] in R to reduce the physical memory usage by mapping some objects in workers to the same shared address spaces, thus eliding copies—unless, of course, a worker modifies the object. This can be loosely seen as mildly clawing back some of the benefits of forking with respect to memory. Note however, that `mori` works on a limited set of R objects, specifically “atomic vector types, lists, and data frames” [6]. See Section for empirical proof that we cannot naively use `mori` to share compiled model methods across workers.

Since OOM is a slightly tricky thing to think about, one user friendly feature we would like to implement, somewhat orthogonal to swapping to `mirai`, is having parallel functions check how much memory they need and warn users if a run might run out of memory. We could do this pretty simply by running our functions across varying input sizes and modelling peak memory usage (probably averaged over a few runs) as a function of the parameters. Of course, if we implement this we would need to test it to see how accurate and conservative our model is.

`mori` and model methods

- [1] A. Vehtari *et al.*, “Loo: Efficient leave-one-out cross-validation and WAIC for bayesian-models.” 2025. Available: <https://mc-stan.org/loo/>
- [2] C. Gao, *Mirai: Minimalist async evaluation framework for r*. 2026. Available: <https://mirai.r-lib.org>
- [3] J. Hester and D. Vaughan, *Bench: High precision timing of r expressions*. 2025. Available: <https://bench.r-lib.org/>
- [4] R Core Team, *R: A language and environment for statistical computing*. Vienna, Austria: R Foundation for Statistical Computing, 2026. Available: <https://www.R-project.org/>
- [5] H. Wickham, *Advanced r*, 2nd ed. Chapman; Hall/CRC, 2019. doi: [10.1201/9781351201315](https://doi.org/10.1201/9781351201315). Available: <https://adv-r.hadley.nz/>
- [6] C. Gao, “Mori: Shared memory for r objects,” Apr. 23, 2026. Available: https://opensource.posit.co/blog/2026-04-23_mori-0-1-0/. [Accessed: May 26, 2026]